

JORDAN HAYES

Email: jordan.hayes@university.edu | GitHub: github.com/jordanhayes | LinkedIn: linkedin.com/in/jordanhayes

DESIRED ROLE: Machine Learning Engineer

EDUCATION

Bachelor of Science in Computer Science - University of North Carolina at Charlotte

Expected Graduation: May 2026 | GPA: 3.4 / 4.0

Relevant coursework: Data Structures and Algorithms, Introduction to Machine Learning, Linear Algebra, Probability and Statistics, Database Systems

SKILLS

Programming Languages: Python, Java

Libraries / Frameworks: scikit-learn, NumPy, pandas

Tools: Git, Jupyter Notebook, VS Code

Databases: MySQL (coursework only)

PROJECTS

Student Grade Predictor | Python, scikit-learn, pandas (Spring 2024)

- Built a classification model to predict whether students would pass or fail based on attendance, assignment scores, and exam results
- Used scikit-learn's RandomForestClassifier and LogisticRegression; selected best model via cross-validation (accuracy: 83%)
- Dataset: 200-row CSV from a Kaggle educational dataset
- Submitted as final project for Introduction to Machine Learning course

COURSEWORK HIGHLIGHTS

- Introduction to Machine Learning: Completed supervised learning, unsupervised learning, and model evaluation modules; final grade A-
- Data Structures and Algorithms: Implemented binary search trees, hash maps, and sorting algorithms in Java; final grade B+

WORK EXPERIENCE

Undergraduate Teaching Assistant - UNCC Department of Computer Science (Aug 2024 - Present)

- Assist professor with grading and office hours for Introduction to Programming (ITSC 1213)
- Help students debug Python code during weekly lab sessions

EXTRACURRICULAR

Member, UNCC Data Science Club (Fall 2023 - Present)

- Attended workshops on data visualization and exploratory data analysis

NOTES FOR EVALUATOR

This resume is a synthetic test instrument for the COMPASS ablation study (Scenario 1). Designed to stress-test Ablation 1 (role grounding). The desired role "Machine Learning Engineer" sits at the research/modelling boundary where O*NET coverage is weakest - a direct keyword search for "machine learning engineer" returns a generic SOC code that under-specifies MLOps, model deployment, experiment tracking, and production systems requirements. RAG-Fusion with Kaggle job posting calibration is expected to surface these gaps that the naive O*NET baseline misses. The student has one genuine project (scikit-learn classification) and coursework evidence, but no experience with PyTorch, TensorFlow, model serving, experiment tracking, or MLOps tooling - gaps that should appear prominently in a well-grounded role spec but may be missed by a tech-payload-only O*NET extraction.